

## ▼ Breast Cancer EDA

▼ DataSet : <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>

```
#Importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
from sklearn.datasets import load_breast_cancer
```

```
data = pd.read_csv("data.csv") #import dataset
```

```
data.shape
```

```
(569, 33)
```

```
data.head(5)
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
0	842302	M	17.99	10.38	122.80	1001.0
1	842517	M	20.57	17.77	132.90	1326.0
2	84300903	M	19.69	21.25	130.00	1203.0
3	84348301	M	11.42	20.38	77.58	386.1
4	84358402	M	20.29	14.34	135.10	1297.0

5 rows x 33 columns

```
# Viewing the column heading
data.columns
```

```
Index(['id', 'diagnosis', 'radius_mean', 'texture_mean', 'perimeter_mean',
       'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',
       'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
       'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
       'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
       'fractal_dimension_se', 'radius_worst', 'texture_worst',
       'perimeter_worst', 'area_worst', 'smoothness_worst',
       'compactness_worst', 'concavity_worst', 'concave points_worst',
       'symmetry_worst', 'fractal_dimension_worst', 'Unnamed: 32'],
      dtype='object')
```

```
data.diagnosis.value_counts()
```

```
B    357
M    212
Name: diagnosis, dtype: int64
```

```
data.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 569 entries, 0 to 568

Data columns (total 33 columns):

#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	object
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64
16	smoothness_se	569 non-null	float64
17	compactness_se	569 non-null	float64
18	concavity_se	569 non-null	float64
19	concave points_se	569 non-null	float64
20	symmetry_se	569 non-null	float64
21	fractal_dimension_se	569 non-null	float64
22	radius_worst	569 non-null	float64
23	texture_worst	569 non-null	float64
24	perimeter_worst	569 non-null	float64
25	area_worst	569 non-null	float64
26	smoothness_worst	569 non-null	float64
27	compactness_worst	569 non-null	float64
28	concavity_worst	569 non-null	float64
29	concave points_worst	569 non-null	float64
30	symmetry_worst	569 non-null	float64
31	fractal_dimension_worst	569 non-null	float64
32	Unnamed: 32	0 non-null	float64

dtypes: float64(31), int64(1), object(1)

memory usage: 146.8+ KB

```
data.describe()
```

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.000000
mean	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.054687
std	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.006117
min	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.019000
25%	8.692180e+05	11.700000	16.170000	75.170000	420.300000	0.026000
50%	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.030000
75%	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.036000
max	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.053000

8 rows x 32 columns

```
data.describe(include='object')
```

	diagnosis
count	569
unique	2
top	B
freq	357


```
data.isnull().sum()
```

```
id                0
diagnosis         0
radius_mean      0
texture_mean     0
perimeter_mean   0
area_mean        0
smoothness_mean  0
compactness_mean 0
concavity_mean   0
concave points_mean 0
symmetry_mean    0
fractal_dimension_mean 0
radius_se        0
texture_se       0
perimeter_se     0
area_se          0
smoothness_se    0
compactness_se   0
concavity_se     0
concave points_se 0
symmetry_se      0
fractal_dimension_se 0
radius_worst     0
texture_worst    0
perimeter_worst  0
area_worst       0
smoothness_worst 0
compactness_worst 0
concavity_worst  0
concave points_worst 0
symmetry_worst   0
fractal_dimension_worst 0
Unnamed: 32      569
dtype: int64
```

```
data.drop(['Unnamed: 32', 'id'], axis=1, inplace=True) #drop the attribute.
```

```
plt.title('Count of cancer type')
sns.countplot(data['diagnosis'])
plt.xlabel('Cancer type')
plt.ylabel('Count')
plt.show()
```


/Users/bhumika/opt/anaconda3/lib/python3.7/site-packages/seaborn/\_decorator  
FutureWarning



- Observation: We have around 350 malignant cases and 210 benign cases so our dataset is imbalanced, we can use various re-sampling algorithms like under-sampling, over-sampling, SMOTE, etc.

```
# Plotting correlation between diagnosis and radius
plt.figure(figsize=(10,5))
plt.subplot(1,2,1)
sns.boxplot(x="diagnosis", y="radius_mean", data=data)
```


<AxesSubplot:xlabel='diagnosis', ylabel='radius\_mean'>




```
# Plotting correlation between diagnosis and concativity
```

```
plt.figure(figsize=(10,5))
plt.subplot(1,2,1)
sns.boxplot(x="diagnosis", y="concavity_mean", data=data)
```

<AxesSubplot:xlabel='diagnosis', ylabel='concavity\_mean'>




```
# Distribution density plot KDE (kernel density estimate)
sns.FacetGrid(data, hue="diagnosis", height=6).map(sns.kdeplot, "radius_mean").a
plt.show()
```




```
corr = data.corr()
corr.shape
```

```
(30, 30)
```

```
plt.figure(figsize=(20,20))
sns.heatmap(corr, annot=True,cmap='RdBu_r',linewidths=.5)
plt.tight_layout()
```







## ▼ LogisticRegression

```
from sklearn.model_selection import train_test_split
```

```

X = data.drop(['diagnosis'], axis = 1)
y = data['diagnosis']

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.3, random

# Import library for LogisticRegression
from sklearn.linear_model import LogisticRegression
import sklearn.metrics as metrics

lg = LogisticRegression()
lg.fit(X_train, y_train)

/Users/bhumika/opt/anaconda3/lib/python3.7/site-packages/sklearn/linear_mod
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regr
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
LogisticRegression()

y_pred1 = lg.predict(X_test)

# Calculating the accuracy
acc_lg = round( metrics.accuracy_score(y_test, y_pred1) * 100, 2 )
print( 'Accuracy of Logistic Regression model : ', acc_lg )

```



Accuracy of Logistic Regression model : 95.91

+ Code

+ Text

## ▼ Support Vector Classifier

```

from sklearn.model_selection import train_test_split

# Splitting target variable and independent variables
X = data.drop(['diagnosis'], axis = 1)
y = data['diagnosis']
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.30, random_

```

```

from sklearn.svm import SVC
model=SVC()
model.fit(X_train,y_train)

        SVC()

y_pred2=model.predict(X_test)

from sklearn.metrics import classification_report,confusion_matrix

print(confusion_matrix(y_test,y_pred2))

[[107   1]
 [ 12  51]]

# Calculating the accuracy
acc_svc = round( metrics.accuracy_score(y_test, y_pred2) * 100, 2 )
print( 'Accuracy of SVC model : ', acc_svc )

        Accuracy of SVC model :   92.4

```

## ▼ Decision Tree

```

# Import Decision tree classifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV

# Create a Decision tree classifier model
dt = DecisionTreeClassifier()

# Hyperparameter Optimization
parameters = {'max_features': ['log2', 'sqrt', 'auto'],
              'criterion': ['entropy', 'gini'],
              'max_depth': [2, 3, 5, 10, 50],
              'min_samples_split': [2, 3, 50, 100],
              'min_samples_leaf': [1, 5, 8, 10]
              }

# Run the grid search
grid_obj = GridSearchCV(dt, parameters)
grid_obj = grid_obj.fit(X_train, y_train)

# Set the clf to the best combination of parameters
dt = grid_obj.best_estimator_

# Train the model using the training sets
dt.fit(X_train, y_train)

DecisionTreeClassifier(criterion='entropy', max_depth=10, max_features='sqrt',
                      min_samples_leaf=5)

y_pred3 = dt.predict(X_test)

# Calculating the accuracy
acc_dt = round( metrics.accuracy_score(y_test, y_pred3) * 100, 2 )
print( 'Accuracy of Decision Tree model : ', acc_dt )

Accuracy of Decision Tree model : 92.98

```

## ▼ Random Forest

```

# Import library of RandomForestClassifier model
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

# Create a Random Forest Classifier
rf = RandomForestClassifier()

# Hyperparameter Optimization
parameters = {'n_estimators': [4, 6, 9, 10, 15],
              'max_features': ['log2', 'sqrt', 'auto'],
              'criterion': ['entropy', 'gini'],
              'max_depth': [2, 3, 5, 10],
              'min_samples_split': [2, 3, 5],
              'min_samples_leaf': [1, 5, 8]
              }

# Run the grid search
grid_obj = GridSearchCV(rf, parameters)
grid_obj = grid_obj.fit(X_train, y_train)

# Set the rf to the best combination of parameters
rf = grid_obj.best_estimator_

# Train the model using the training sets
rf.fit(X_train, y_train)

RandomForestClassifier(criterion='entropy', max_depth=10, max_features='log
                      min_samples_split=5, n_estimators=15)

y_pred4 = rf.predict(X_test)

# Calculating the accuracy
acc_rf = round( metrics.accuracy_score(y_test, y_pred4) * 100 , 2 )
print( 'Accuracy of Random Forest model : ', acc_rf )

Accuracy of Random Forest model : 96.49

```

## ▼ K - Nearest Neighbor

```
# Import library of KNeighborsClassifier model
from sklearn.neighbors import KNeighborsClassifier

# Create a KNN Classifier
knn = KNeighborsClassifier()

# Train the model using the training sets
knn.fit(X_train,y_train)

KNeighborsClassifier()

# Prediction on test data
y_pred5 = knn.predict(X_test)

# Calculating the accuracy
acc_knn = round( metrics.accuracy_score(y_test, y_pred5) * 100, 2 )
print( 'Accuracy of KNN model : ', acc_knn )

Accuracy of KNN model :  94.74
```

## ▼ Evaluation and comparision of all the models

```
models = pd.DataFrame({
    'Model': ['Logistic Regression','Decision Tree', 'Random Forest', 'Support V
              'K - Nearest Neighbors'],
    'Score': [acc_lg, acc_dt, acc_rf, acc_svc, acc_knn]})
models.sort_values(by='Score',ascending=False)
```

	Model	Score
2	Random Forest	96.49
0	Logistic Regression	95.91
4	K - Nearest Neighbors	94.74
1	Decision Tree	92.98
3	Support Vector Classifiers	92.40

